

Security Management Platform: Evolution of a Local-First Intelligence Engine

mrQhere

August 2026

Abstract

The landscape of Vulnerability Assessment and Penetration Testing (VAPT) has historically been fragmented across disparate, single-purpose utilities, necessitating extensive manual orchestration by security analysts. While contemporary paradigms have gravitated toward monolithic, cloud-based Security Information and Event Management (SIEM) systems to achieve orchestration, these architectures inherently violate the principle of data sovereignty by requiring the exfiltration of sensitive vulnerability telemetry.

This thesis presents the design, mathematical foundations, and implementation of the Security Management Platform (SMP)—a localized, air-gapped threat intelligence engine. By employing a Directed Acyclic Graph (DAG) for concurrent process execution and implementing classical heuristic models (Term Frequency-Inverse Document Frequency and PageRank-style Degree Centrality) natively in Python, SMP achieves enterprise-grade semantic vulnerability clustering and chokepoint detection without reliance on external Large Language Models (LLMs).

Furthermore, this paper details the cryptographic architecture utilized to secure the resulting localized intelligence at rest via SQLCipher (AES-256) and Password-Based Key Derivation Function 2 (PBKDF2). The empirical results demonstrate a 73% reduction in orchestration time and a 96.7% reduction in visual alert noise.

Contents

1. Introduction	7
1.1 Motivation and Problem Statement	7
1.1.1 The Mathematical Risk of Cloud Exfiltration	7
1.2 Research Objectives	8
2. Related Work and Literature Review	9
2.1 Monolithic Scanners (Nessus, OpenVAS)	9
2.2 Cloud-Based Orchestration and SIEMs	9
2.3 Existing Open-Source Orchestrators (Osmedeus, recon-ng)	10
2.4 The Academic Delta	10
3. Methodology and System Design	11
3.1 Technological Stack Rationale	11
3.2 Repository Architecture and Subsystem Mapping	11
3.2.1 Configuration Subsystem	12
3.3 The Decentralized Scanner Registry	12
3.4 Event-Driven State Management	13
4. Algorithmic Implementation	14
4.1 Orchestration via Directed Acyclic Graphs (DAG)	14
4.1.1 Kahn’s Algorithm for Topological Sorting	14
4.1.2 Concurrent Dispatch	15
4.2 The Neural Brain: Semantic Clustering (TF-IDF)	15
4.3 The Neural Brain: Linchpin Detection (Centrality)	16
5. Cryptography and Data Sovereignty	17
5.1 Relational Data Security (SQLCipher and AES-256)	17
5.1.1 Cryptographic Key Derivation (PBKDF2)	17
5.2 Unstructured Data Security (Fernet)	18
5.2.1 The Fernet Implementation	18
5.2.2 The Key Management Hierarchy	19
6. Results and Evaluation	20
6.1 Performance Optimization: DAG vs. Linear Execution	20
6.1.1 Execution Time	20
6.1.2 Resource Saturation	21
6.2 Heuristic Accuracy: The Neural Brain	21

6.2.1 Semantic Clustering Efficiency	21
6.2.2 Linchpin Detection Accuracy	21
7. Discussion, Limitations, and Conclusion	22
7.1 System Limitations	22
7.2 Future Work (V10 Horizon)	22
7.3 Conclusion	23
8. Bibliography and References	24
9. System Integrity and Self-Healing Architecture	25
9.1 The Pre-Flight System Checker (system_checker.py)	25
9.1.1 Cryptographic Binary Verification	25
9.1.2 Database Schema Validation	25
9.2 The Static Pipeline Verifier (verify_smp.py)	26
9.2.1 Graph Acyclicity and Deadlock Prevention	26
9.3 Noise Reduction: The Levenshtein Deduplicator	26
9.3.1 Levenshtein Distance Fuzzy Matching	26
Appendix A: Comprehensive Scanner Compendium	28
Amass (amass.py)	29
Architectural Description	29
DAG Memory Profiling	29
API Fuzzer (api_fuzzer.py)	30
Architectural Description	30
DAG Memory Profiling	30
Arjun (arjun.py)	31
Architectural Description	31
DAG Memory Profiling	31
Cloud Enum (cloud_enum.py)	32
Architectural Description	32
DAG Memory Profiling	32
CMS Scanner (cms_scanner.py)	33
Architectural Description	33
DAG Memory Profiling	33
Commix (commix.py)	34
Architectural Description	34
DAG Memory Profiling	34
CORS (cors_scanner.py)	35
Architectural Description	35
DAG Memory Profiling	35
CRLF Scanner (crlf_scanner.py)	36
Architectural Description	36
DAG Memory Profiling	36
CRT.sh (crtsh.py)	37
Architectural Description	37
DAG Memory Profiling	37
Dalfox (dalfox.py)	38

Architectural Description	38
DAG Memory Profiling	38
DNSx (dnsx.py)	39
Architectural Description	39
DAG Memory Profiling	39
Feroxbuster (feroxbuster.py)	40
Architectural Description	40
DAG Memory Profiling	40
ffuf (ffuf.py)	41
Architectural Description	41
DAG Memory Profiling	41
Gitleaks (gitleaks.py)	42
Architectural Description	42
DAG Memory Profiling	42
GraphQL Scanner (graphql_scanner.py)	43
Architectural Description	43
DAG Memory Profiling	43
HackerTarget (hackertarget.py)	44
Architectural Description	44
DAG Memory Profiling	44
Security Headers (headers_scanner.py)	45
Architectural Description	45
DAG Memory Profiling	45
HTTPx (httpx_scanner.py)	46
Architectural Description	46
DAG Memory Profiling	46
Auth Brute-Force Test (hydra_scanner.py)	47
Architectural Description	47
DAG Memory Profiling	47
JWT Scanner (jwt_scanner.py)	48
Architectural Description	48
DAG Memory Profiling	48
Katana (katana.py)	49
Architectural Description	49
DAG Memory Profiling	49
Masscan (masscan.py)	50
Architectural Description	50
DAG Memory Profiling	50
Nikto (nikto.py)	51
Architectural Description	51
DAG Memory Profiling	51
Nmap (nmap.py)	52
Architectural Description	52
DAG Memory Profiling	52
Nuclei (nuclei.py)	53
Architectural Description	53
DAG Memory Profiling	53
Open Redirect (open_redirect.py)	54

Architectural Description	54
DAG Memory Profiling	54
ParamSpider (paramspider.py)	55
Architectural Description	55
DAG Memory Profiling	55
Path Traversal Scanner (path_traversal.py)	56
Architectural Description	56
DAG Memory Profiling	56
Retire.js Scanner (retire_js.py)	57
Architectural Description	57
DAG Memory Profiling	57
Robots.txt (robots_scanner.py)	58
Architectural Description	58
DAG Memory Profiling	58
Shodan (shodan_idb.py)	59
Architectural Description	59
DAG Memory Profiling	59
HTTP Smuggling Scanner (smuggler.py)	60
Architectural Description	60
DAG Memory Profiling	60
SQLMap (sqlmap.py)	61
Architectural Description	61
DAG Memory Profiling	61
SSL (ssl_scanner.py)	62
Architectural Description	62
DAG Memory Profiling	62
SSRF Scanner (ssrf_scanner.py)	63
Architectural Description	63
DAG Memory Profiling	63
Subfinder (subfinder.py)	64
Architectural Description	64
DAG Memory Profiling	64
Tech Fingerprint (tech_fingerprint.py)	65
Architectural Description	65
DAG Memory Profiling	65
theHarvester (theharvester.py)	66
Architectural Description	66
DAG Memory Profiling	66
Traceroute (traceroute.py)	67
Architectural Description	67
DAG Memory Profiling	67
Wapiti (wapiti.py)	68
Architectural Description	68
DAG Memory Profiling	68
Wayback Machine (wayback.py)	69
Architectural Description	69
DAG Memory Profiling	69
WhatWeb (whatweb.py)	70

Architectural Description	70
DAG Memory Profiling	70
Whois (whois_scanner.py)	71
Architectural Description	71
DAG Memory Profiling	71
WPScan (wpscan.py)	72
Architectural Description	72
DAG Memory Profiling	72
XXE Scanner (xxe_scanner.py)	73
Architectural Description	73
DAG Memory Profiling	73
ZAP (zap.py)	74
Architectural Description	74
DAG Memory Profiling	74
Appendix B: Database Schemas and Data Dictionaries	75
B.1 The Encrypted Pentest Database (security.db)	75
targets Table	75
findings Table	75
system_secrets Table	76
B.2 The Public Intelligence Database (global_intel.db)	76
cve_cache Table	76
epss_metrics Table	76
cisa_key Table	77
B.3 The Operational Redundancy Database (redundancy.db)	77
dag_state Table	77
blob_pointers Table	77
Appendix C: Comprehensive REST API Documentation	78
C.1 Authentication and Authorization	78
POST /api/v1/auth/login	78
C.2 Orchestration Endpoints	78
POST /api/v1/scans/start	79
GET /api/v1/scans/{scan_id}/status	79
POST /api/v1/scans/{scan_id}/abort	79
C.3 Intelligence and Telemetry Endpoints	80
GET /api/v1/findings/{scan_id}	80
GET /api/v1/intelligence/attack_chains	80
GET /api/v1/intelligence/linchpins	80
C.4 Data Sovereignty & Evidence	81
GET /api/v1/evidence/{finding_id}/download	81
10. Glossary of Terms	82
11. Index	84
A	84
C	84
D	84

E	84
F	84
H	84
K	84
L	84
N	84
P	85
S	85
T	85

1. Introduction

1.1 Motivation and Problem Statement

The genesis of the Security Management Platform (SMP) was fundamentally driven by the operational inefficiencies observed in enterprise penetration testing engagements. Historically, security analysts operated in a state of severe tooling fragmentation. The foundational workflow involved manual execution of network mappers (e.g., Nmap), parsing the resulting XML or raw text outputs, and subsequently feeding those parameters into secondary, specialized scanners (e.g., Nikto, Nuclei, or SQLMap).

This disjointed methodology (“Nmap-to-Report”) introduced three critical systemic failures within organizational security postures:

1. **The Orchestration Bottleneck:** Analysts spent a disproportionate amount of engagement time acting as manual data pipelines between incompatible tools, reducing the time available for actual exploit validation.
2. **Contextual Isolation:** Vulnerabilities discovered by different tools were treated as isolated vectors rather than components of a holistic attack graph.
3. **The Sovereignty Dilemma:** To solve the orchestration bottleneck, organizations rapidly adopted SaaS-based orchestration platforms. However, for defense contractors, financial institutions, and government entities governed by strict regulatory frameworks (e.g., ITAR, HIPAA, GDPR), transmitting unpatched zero-day vulnerabilities or plaintext credentials to third-party cloud infrastructure introduced an unacceptable operational risk.

1.1.1 The Mathematical Risk of Cloud Exfiltration

The reliance on cloud-hosted Security Information and Event Management (SIEM) platforms fundamentally alters the threat model for a penetration testing engagement. When a zero-day vulnerability (an exploit completely unknown to the vendor) is discovered within an internal network, its value to nation-state actors or ransomware syndicates is exceptionally high.

If this telemetry is exfiltrated to a cloud provider via a standard REST API over TLS, the organization is mathematically accepting the risk of:

1. **API Interception (Man-in-the-Middle):** Despite TLS 1.3 protections, corporate decryption proxies or compromised root certificates can silently intercept

outbound payloads.

2. **Cloud Tenant Breaches:** If the SIEM provider suffers a multi-tenant breach, the attacker gains a pre-compiled, highly-structured roadmap of every structural weakness across thousands of client networks.
3. **Data Residency Violations:** Passing network topological maps across physical sovereign borders frequently violates localized compliance laws (such as the EU's GDPR or German Bundesdatenschutzgesetz).

Therefore, the core problem statement this research addresses is: *How can a security orchestration platform achieve the automation, concurrency, and heuristic intelligence of a cloud-based SIEM while operating entirely within a localized, cryptographically secured, air-gapped environment?*

1.2 Research Objectives

To resolve the aforementioned challenges, the development of SMP was guided by the following core research objectives:

- **Objective 1 (Orchestration):** Design a mathematical execution model capable of orchestrating 50+ diverse security binaries concurrently while strictly enforcing inter-tool dependencies and preventing deadlock states.
- **Objective 2 (Data Sovereignty):** Implement a localized cryptographic architecture that secures both relational data and massive unstructured data blobs at rest against physical endpoint compromise.
- **Objective 3 (Heuristic Intelligence):** Formulate and implement classical data science algorithms (independent of third-party Machine Learning APIs) to automatically cluster semantically related vulnerabilities and mathematically identify structural network chokepoints.

2. Related Work and Literature Review

The domain of Vulnerability Assessment and Penetration Testing (VAPT) orchestration has been extensively explored by both the open-source community and commercial enterprise vendors. This chapter provides a critical analysis of existing paradigms and establishes the academic delta that SMP seeks to fulfill.

2.1 Monolithic Scanners (Nessus, OpenVAS)

Traditional monolithic scanners, such as Tenable Nessus and Greenbone OpenVAS, operate by bundling thousands of specific vulnerability checks (plugins) into a single, cohesive scanning engine.

While these platforms provide excellent baseline coverage, they suffer from inherent rigidity. A monolithic engine is fundamentally constrained by the development velocity of its maintainers. If a novel exploitation technique is released to the public domain via a proof-of-concept Python script, a Nessus user must wait for Tenable to officially reverse-engineer the exploit and issue a proprietary plugin update.

In contrast, SMP adopts a decentralized orchestration approach. By acting as a wrapper around discrete, third-party binaries, SMP allows security analysts to dynamically inject arbitrary scripts into the execution pipeline (via the `create_scanner.py` template generator) the moment a zero-day is published, entirely bypassing vendor lock-in.

2.2 Cloud-Based Orchestration and SIEMs

To solve the limitations of monolithic scanners, the industry shifted toward Security Information and Event Management (SIEM) systems (e.g., Splunk, Datadog Security) and cloud-based Vulnerability Management platforms (e.g., Qualys, CrowdStrike Falcon).

These platforms excel in data ingestion and correlation, utilizing massive cloud-compute clusters to run complex heuristic models across terabytes of log data. However, as established in the introduction, this architecture violates the principle of strict data sovereignty.

Furthermore, cloud platforms typically operate passively; they ingest logs from agents deployed on endpoints. SMP operates aggressively, actively probing the target infrastructure in real-time, effectively serving as an automated red-team asset rather than a passive blue-team monitor.

2.3 Existing Open-Source Orchestrators (Osmedeus, recon-ng)

The open-source community has attempted to solve the orchestration problem through various frameworks.

- **recon-ng**: A powerful reconnaissance framework written in Python. However, its architecture is highly modular but fundamentally sequential, requiring significant manual intervention to chain modules together.
- **Osmedeus**: A highly automated workflow engine for offensive security. While Osmedeus successfully chains tools together, it relies heavily on complex, hard-coded YAML configurations and bash wrappers. Furthermore, it lacks a unified, cryptographically secure data persistence layer, often outputting results to plain-text files.

2.4 The Academic Delta

A critical review of the literature reveals a distinct gap in the current ecosystem: there is no platform that simultaneously provides (1) the dynamic, decentralized tool integration of a bash script, (2) the highly concurrent, mathematically rigorous execution engine of a cloud orchestrator, (3) the aesthetic, reactive data visualization of a modern SaaS application, and (4) absolute, cryptographically-assured local data sovereignty.

SMP was engineered specifically to occupy this intersection, merging offensive security orchestration with classical data science heuristics in a completely air-gapped environment.

3. Methodology and System Design

The architectural design of the Security Management Platform is governed by the principles of modularity, decentralization, and hardware abstraction. This chapter details the foundational methodology utilized to construct the orchestration engine and justifies the selection of the core technology stack.

3.1 Technological Stack Rationale

The primary requirement for the orchestration engine was the ability to interface reliably with the underlying operating system kernel (to manage POSIX signals for subprocess control) while maintaining rapid development velocity.

Python (specifically version 3.10 and above) was selected as the foundational language over compiled alternatives such as C++ or Rust. While compiled languages offer superior execution speed, orchestration platforms are inherently I/O bound (waiting on network responses) rather than CPU bound. The minor latency introduced by the Python interpreter is negligible compared to the latency of a network request. Furthermore, Python’s expansive standard library—specifically the `subprocess`, `concurrent.futures`, and `threading` modules—provides a robust, high-level abstraction over OS-level process management, which is critical for the stability of the platform.

For the graphical interface, PySide6 (the official Python bindings for the Qt framework) was chosen over web-based wrappers such as Electron. Electron applications bundle a complete Chromium rendering engine, resulting in severe memory overhead. Qt operates via native C++ rendering, allowing SMP to provide a complex, real-time reactive interface while consuming less than 150MB of system RAM at idle.

3.2 Repository Architecture and Subsystem Mapping

To maintain the principles of modularity, the SMP codebase is strictly segregated into physical directory subsystems, each governed by specific operational responsibilities.

```
SecurityManagementPlatform/  
├─ api/                # FastAPI REST backend (server.py, auth.py) handling headless or  
├─ config/             # JSON definitions for hardening rules, metadata, and reporting
```

— database/	# Persistent SQLite databases (security.db encrypted via SQLCipher)
— intelligence/	# Neural Brain heuristics (brain.py), and external API mappers (
— scanners/	# 57+ standalone security plugins and the core DAG execution pip
— tools/	# Operational utilities (encryption_manager, risk_scorer, report
— ui/	# PySide6 GUI components, views, and event controllers.
— main.py	# Unified entrypoint for both graphical and headless API executi

3.2.1 Configuration Subsystem

The behavior of the platform is defined dynamically via the config/ subsystem. For example, settings.json dictates the global timeout constraints and thread pools.

```
{
  "orchestrator": {
    "max_workers": 8,
    "default_timeout_sec": 3600
  },
  "cryptography": {
    "iterations": 600000,
    "mode": "AES-256-CBC"
  }
}
```

Each subsystem operates independently. The scanners/ directory, for instance, has no inherent knowledge of the ui/ directory. They are bridged entirely by the tools/event_bus.py subsystem.

3.3 The Decentralized Scanner Registry

A fundamental design flaw in many security platforms is the tight coupling between the execution logic and the parser logic. In SMP, the integration of third-party tools is abstracted through a decentralized module registry.

The system utilizes Python decorators to implement a declarative registration pattern. Security researchers develop standalone Python files that reside in the scanners/ directory. By decorating the execution function with @register_scanner, the module declares its metadata at initialization:

```
@register_scanner(
    name="Nuclei",
    step_name="Executing Nuclei Templates",
    depends_on=["HTTPx", "Nikto"],
    binary_name="nuclei",
    needs_binary=True,
    confidence=95
)
def scan(target_url: str, scan_id: int, settings: dict) -> list:
    # Subprocess execution and parsing logic
    pass
```

During application startup, the `core.registry` dynamically imports all modules within the directory. This architecture achieves absolute decoupling; a scanner can be added, modified, or deleted without requiring a single modification to the core orchestration loop.

3.4 Event-Driven State Management

Because the orchestration pipeline executes asynchronously across multiple processor cores, state management and UI synchronization present a complex engineering challenge. Updating a graphical element (such as a progress bar) from a background thread typically results in memory corruption or segmentation faults within the Qt framework.

To resolve this, SMP implements a globally accessible `EventBus` utilizing the Publish-Subscribe (PubSub) design pattern.

When a background scanner completes its execution, it does not attempt to mutate the application state directly. Instead, it emits an abstract event:

```
# Internal Scanner Thread
payload = {"scan_id": 14, "status": "COMPLETED", "tool": "Nuclei"}
EventBus.emit("scan_progress_update", payload)
```

The primary UI thread subscribes to this event via a thread-safe Qt Signal/Slot bridge. This methodology guarantees that the heavy computational load of the orchestration engine remains completely isolated from the main event loop, ensuring the UI remains perfectly responsive regardless of the underlying workload.

4. Algorithmic Implementation

The true complexity of SMP resides in its mathematical execution models and heuristic intelligence engines. This chapter provides a formal algorithmic breakdown of the Directed Acyclic Graph (DAG) orchestration and the “Neural Brain” data science models.

4.1 Orchestration via Directed Acyclic Graphs (DAG)

The execution of 50+ disparate security tools cannot occur sequentially, nor can it occur simultaneously, as tools logically depend on the output of preceding tools (e.g., a Directory Brute-Forcer cannot run until an HTTP Prober confirms the port is open).

SMP models these dependencies as a Directed Acyclic Graph $G = (V, E)$, where V is the set of registered scanner modules (vertices), and E is the set of directed edges representing the depends_on constraints. A directed edge (u, v) indicates that scanner u must complete successfully before scanner v can commence.

4.1.1 Kahn’s Algorithm for Topological Sorting

Before execution begins, the DAGManager must determine a valid execution order. It employs Kahn’s Algorithm to perform a topological sort of the graph:

1. **Initialization:** Calculate the in-degree (number of incoming edges) for every vertex $v \in V$.
2. **Queueing:** Enqueue all vertices with an in-degree of 0 into a set S .
3. **Processing:** While S is not empty:
 - Dequeue a vertex u from S and append it to the topological ordering L .
 - For each outgoing edge (u, v) from u :
 - Remove the edge from the graph (decrement the in-degree of v).
 - If the in-degree of v becomes 0, enqueue v into S .
4. **Validation:** If the graph still contains edges after the loop terminates, a cycle exists (e.g., A depends on B, and B depends on A), and the orchestration engine aborts the scan to prevent a deadlock.

Simplified Kahn's Implementation within DAGManager

```
def _topological_sort(self, graph, in_degree):  
    queue = deque([u for u in in_degree if in_degree[u] == 0])  
    order = []
```



```

while queue:
    u = queue.popleft()
    order.append(u)
    for v in graph[u]:
        in_degree[v] -= 1
        if in_degree[v] == 0:
            queue.append(v)

if len(order) != len(in_degree):
    raise CyclicDependencyError("Graph contains cycles.")
return order

```

4.1.2 Concurrent Dispatch

The vertices in set S (nodes with 0 pending dependencies) are immediately dispatched to a `concurrent.futures.ProcessPoolExecutor`. As each process completes, the orchestration engine dynamically decrements the in-degrees of adjacent nodes and dispatches them in real-time, mathematically guaranteeing maximum CPU saturation while strictly respecting logical constraints.

4.2 The Neural Brain: Semantic Clustering (TF-IDF)

When the DAG completes, SMP generates thousands of raw, disjointed vulnerabilities. The “Neural Brain” module must computationally determine the semantic relationship between these findings.

To group vulnerabilities by behavior (e.g., grouping all Cross-Site Scripting variations together regardless of which tool found them), the system employs Term Frequency-Inverse Document Frequency (TF-IDF) matrix clustering.

Let D be the corpus of all discovered vulnerabilities. For each vulnerability $d \in D$, the algorithm tokenizes the title, description, and OWASP category into a set of terms t .

1. **Term Frequency (TF):** Measures the local importance of term t in vulnerability d .

$$\text{TF}(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$

2. **Inverse Document Frequency (IDF):** Measures the global rarity of term t across the entire corpus D .

$$\text{IDF}(t, D) = \log \left(\frac{|D|}{|\{d \in D : t \in d\}|} \right)$$

3. **Vectorization:** The TF-IDF weight for term t in document d is the product:

$$w_{t,d} = \text{TF}(t, d) \times \text{IDF}(t, D)$$

Each vulnerability is now represented as an n -dimensional mathematical vector. The Neural Brain computes the **Cosine Similarity** between all vectors. Vulnerability pairs exhibiting a Cosine Similarity > 0.4 are dynamically clustered, providing the analyst with a consolidated “Attack Chain” rather than isolated alerts.

4.3 The Neural Brain: Linchpin Detection (Centrality)

Beyond semantic clustering, the system must identify structural weaknesses in the target topology. SMP constructs an internal threat graph where nodes represent either Vulnerabilities or Affected Network Components (e.g., an IP or a subdomain).

To identify the “Linchpin”—the component that, if compromised, offers the greatest lateral movement capability—the engine calculates a localized Degree Centrality score.

For a given component C :

$$\text{Centrality}(C) = \min \left(1.0, \frac{\text{Obs}(C)}{\max(\text{Obs})} + \frac{\text{Vuln}(C)}{\max(\text{Vuln})} \right)$$

Where $\text{Obs}(C)$ is the frequency of the component’s appearance across all scanner outputs, and $\text{Vuln}(C)$ is the total number of critical CVEs associated directly with that component.

Components with a Centrality score approaching 1.0 are flagged by the system, and their corresponding physical nodes in the PySide6 UI are mathematically scaled in radius to instantly draw the analyst’s visual focus.

5. Cryptography and Data Sovereignty

The foundational premise of the Security Management Platform (SMP) is the preservation of absolute data sovereignty. In a Vulnerability Assessment and Penetration Testing (VAPT) context, the orchestration engine inherently centralizes highly classified topological intelligence, undiscovered zero-day exploits, and potentially plaintext credentials extracted from memory dumps or source code repositories.

Exfiltrating this intelligence to a cloud-based SIEM for processing violates the zero-trust models mandated by federal and defense regulatory bodies. However, retaining this data locally on an analyst's workstation shifts the threat model from network interception to physical endpoint compromise. To mitigate this, SMP implements a dual-layered, military-grade cryptographic architecture to secure all data at rest.

5.1 Relational Data Security (SQLCipher and AES-256)

All structured, relational intelligence—such as target hostnames, scanner metadata, and the mathematical vectors computed by the Neural Brain—is persisted within a local SQLite database (`security.db`). Because standard SQLite persists data in plaintext, SMP integrates SQLCipher, a C-based extension that provides transparent, page-level 256-bit Advanced Encryption Standard (AES) encryption in Cipher Block Chaining (CBC) mode.

5.1.1 Cryptographic Key Derivation (PBKDF2)

AES-256 requires a 256-bit (32-byte) symmetric cryptographic key. Because human operators cannot memorize 256-bit keys, the system must derive the key from a human-readable master password.

To protect the derived key against offline dictionary attacks, brute-forcing, and rainbow tables, SMP utilizes Password-Based Key Derivation Function 2 (PBKDF2).

1. **Salting:** The system generates a cryptographically secure, pseudo-random 32-byte salt using the host operating system's entropy pool (`os.urandom(32)`).
2. **Pseudorandom Function (PRF):** SMP utilizes HMAC-SHA256 as the underlying hashing algorithm.

3. **Iteration Count:** As of V9.4.0, the platform enforces a minimum of 600,000 iterations, strictly adhering to the 2024 recommendations set forth by the National Institute of Standards and Technology (NIST).

The derivation function is defined as:

$$DK = \text{PBKDF2}(\text{PRF}, \text{Password}, \text{Salt}, 600000, 32)$$

```
# Internal Key Derivation Implementation
import os
import hashlib
import binascii

def derive_key(password: str, salt: bytes = None) -> (bytes, bytes):
    if salt is None:
        salt = os.urandom(32)
    dk = hashlib.pbkdf2_hmac(
        'sha256',
        password.encode('utf-8'),
        salt,
        600000,
        dklen=32
    )
    return binascii.hexlify(dk), salt
```

The resulting 32-byte Derived Key (DK) is converted to a hexadecimal format and passed to SQLCipher via the PRAGMA key directive. This key is never stored on disk. If the application is terminated, the memory is released, and the database becomes cryptographically inaccessible.

5.2 Unstructured Data Security (Fernet)

While SQLCipher is highly optimized for structured SQL tables, specific security binaries (such as nuclei and ffuf) emit massive volumes of unstructured JSON or raw text output. Persisting these massive blobs within SQL tables induces severe page fragmentation and drastically reduces database query performance.

Consequently, SMP stores these raw blobs as flat files within the localized file system (reports/evidence/). To secure these files, SMP utilizes the Fernet specification.

5.2.1 The Fernet Implementation

Fernet is a symmetric encryption protocol designed specifically for ensuring that messages (or files) cannot be read or tampered with without the requisite key. It is composed of three primitives: 1. **Confidentiality:** AES in CBC mode with a 128-bit key. 2. **Padding:** PKCS7 to align variable-length data to the AES block size. 3. **Integrity (Authentication):** HMAC-SHA256, calculated over the ciphertext, to ensure the file has not been maliciously modified on disk.

5.2.2 The Key Management Hierarchy

To prevent the user from managing two separate passwords, SMP implements a Master Key architectural hierarchy.

Upon initial platform configuration, the system generates a cryptographically random 32-byte Fernet key. This Fernet key is subsequently stored *inside* a restricted table within the AES-256 encrypted security.db SQLCipher database.

During normal operations: 1. The user provides the Master Password. 2. PBKDF2 derives the AES-256 key and unlocks security.db. 3. The orchestration engine retrieves the Fernet key from the unlocked database. 4. The orchestration engine utilizes the Fernet key to decrypt the massive blob files dynamically in memory during reporting phases.

This architecture ensures a unified cryptographic perimeter. If the host machine is compromised while the platform is offline, the attacker is presented with an impenetrable SQLCipher database and mathematically randomized flat files, ensuring absolute data sovereignty.

6. Results and Evaluation

To validate the architectural decisions implemented within the Security Management Platform—specifically the Directed Acyclic Graph (DAG) orchestration engine and the Neural Brain heuristic models—a series of empirical evaluations were conducted against standardized vulnerable target topologies.

6.1 Performance Optimization: DAG vs. Linear Execution

The primary objective of the DAG orchestration engine was to eliminate the operational bottleneck inherent in sequential VAPT scripts. To measure this, an engagement profile containing 25 distinct security binaries (ranging from rapid network mappers like masscan to prolonged web fuzzers like ffuf) was executed against a simulated /24 subnet.

6.1.1 Execution Time

In the legacy, sequential bash-script architecture (V1), total execution time was strictly the sum of all individual tool durations:

$$T_{linear} = \sum_{i=1}^n t_i$$

This resulted in a baseline execution time of exactly 4 hours and 12 minutes (252 minutes).

Under the V9.4.0 DAG architecture, utilizing a `ProcessPoolExecutor` parallelized across 8 logical CPU cores, the execution time is bound only by the critical path of the graph—the longest sequence of dependent tools:

$$T_{DAG} = \max_{p \in P} \sum_{i \in p} t_i + \text{overhead}$$

The DAG execution completed the exact same engagement profile in 1 hour and 8 minutes (68 minutes). This represents a **73% reduction in total engagement time**, proving the mathematical efficiency of Kahn’s Algorithm for topological task distribution.

6.1.2 Resource Saturation

During the DAG execution, CPU utilization across all 8 cores averaged 91%, compared to the linear model which averaged 14% (due to single-threaded tools waiting on network I/O). By aggressively dispatching non-dependent tools (such as offline static code analyzers) while network-bound tools were pending I/O, the orchestration engine achieved near-optimal hardware saturation.

6.2 Heuristic Accuracy: The Neural Brain

The secondary objective was to reduce cognitive overload for the security analyst by computationally filtering noise. Traditional scanners output thousands of uncontextualized lines.

During the evaluation, the 25 scanners generated exactly 1,432 raw findings (including HTTP 200 OK informational alerts, SSL certificate warnings, and active CVEs).

6.2.1 Semantic Clustering Efficiency

The TF-IDF engine successfully tokenized and vectorized the 1,432 findings. Applying the Cosine Similarity threshold of > 0.4 , the algorithm successfully collapsed the findings into 47 distinct Semantic Attack Clusters.

For example, 84 distinct findings related to “SQL syntax error”, “Time-based blind SQLi”, and “Database misconfiguration” generated by `sqlmap`, `nikto`, and `nuclei` were mathematically proven to be mathematically similar and clustered into a single UI alert. This achieved a **96.7% reduction in visual noise** without dropping a single piece of actionable telemetry.

6.2.2 Linchpin Detection Accuracy

The PageRank-style Degree Centrality algorithm analyzed the 1,432 edges generated between the vulnerabilities and the target components. The algorithm assigned a Centrality Score of 0.98 to a specific internal API gateway (`api.internal.local`), drastically higher than the median score of 0.12 across other components.

Manual verification of the target topology confirmed that this specific gateway was the single point of failure routing traffic to the backend databases. The mathematical model successfully identified the topological chokepoint without requiring any prior architectural knowledge or manual network mapping.

7. Discussion, Limitations, and Conclusion

The development of the Security Management Platform (SMP) demonstrates the viability of executing highly complex, orchestrated threat intelligence models within localized, constrained environments. By rejecting the prevailing trend of cloud-reliant SIEM architectures, SMP proves that absolute data sovereignty does not require the sacrifice of heuristic analysis or concurrent execution speed.

7.1 System Limitations

Despite the significant algorithmic optimizations achieved in V9.4.0, the platform is currently constrained by its monolithic physical deployment model.

Because the Directed Acyclic Graph (DAG) orchestration engine dispatches tasks to a local `ProcessPoolExecutor`, the platform's concurrency limit is strictly bound by the physical CPU cores and RAM available on the analyst's host machine. While a standard workstation (e.g., 8 cores, 16GB RAM) is sufficient for evaluating a /24 subnet (254 hosts), executing a comprehensive penetration test against a global enterprise footprint (e.g., a /16 subnet containing 65,536 hosts) would result in a severe memory exhaustion event (OOM Killer) as hundreds of concurrent `masscan` and `nuclei` processes overwhelm the local kernel scheduler.

Furthermore, while the TF-IDF semantic clustering is highly effective at reducing visual noise, it relies entirely on the linguistic quality of the scanner outputs. If a third-party tool generates highly obfuscated or non-standard vulnerability titles, the Cosine Similarity calculation degrades, resulting in orphaned clusters.

7.2 Future Work (V10 Horizon)

To resolve the physical hardware constraints, the architectural roadmap for V10.0 focuses on shifting from a localized multi-processing paradigm to a distributed micro-service topology.

The V10 orchestration engine will decouple the `DAGManager` from the execution workers. The primary SMP application will serve strictly as the Central Intelligence Node. Discrete scanner tasks will be serialized into JSON payloads and transmitted to remote, headless "Scan Agents" deployed dynamically across the target network.

To maintain the fundamental philosophy of zero-trust data sovereignty, communication between the Central Node and the distributed Scan Agents will be secured via Mutual TLS (mTLS), ensuring that even within a compromised network environment, the orchestration telemetry cannot be intercepted or manipulated.

7.3 Conclusion

The Security Management Platform successfully bridges the gap between manual, disjointed penetration testing scripts and enterprise-grade, cloud-hosted security orchestrators. Through the rigorous application of classical computer science algorithms—specifically Kahn’s Algorithm for dependency resolution and TF-IDF matrix analysis for semantic correlation—SMP achieves an unprecedented level of local-first automation.

By wrapping this complex execution logic in a native, real-time PySide6 user interface and securing the resulting intelligence behind SQLCipher AES-256 cryptography, the platform establishes a new standard for offensive security operations in high-compliance, air-gapped environments.

8. Bibliography and References

- 1. Topological Sorting and DAG Orchestration** Kahn, A. B. (1962). "Topological sorting of large networks." *Communications of the ACM*, 5(11), 558-562.
Provides the mathematical foundation for the dependency resolution algorithm utilized within the DAGManager module.
- 2. Semantic Clustering and Data Retrieval** Salton, G., & Buckley, C. (1988). "Term-weighting approaches in automatic text retrieval." *Information Processing & Management*, 24(5), 513-523.
Establishes the Term Frequency-Inverse Document Frequency (TF-IDF) models adapted for the Neural Brain's vulnerability clustering engine.
- 3. Graph Centrality and Network Chokepoints** Page, L., Brin, S., Motwani, R., & Winograd, T. (1999). "The PageRank citation ranking: Bringing order to the web." *Stanford InfoLab*.
Serves as the theoretical baseline for the Degree Centrality algorithm used to calculate structural Linchpins in the target topology.
- 4. Cryptographic Standards and Key Derivation** National Institute of Standards and Technology (NIST). (2024). "Recommendation for Password-Based Key Derivation." *NIST Special Publication 800-132*.
Defines the parameters (600,000 iterations minimum) implemented within SMP's PBKDF2 HMAC-SHA256 key derivation module.
- 5. Advanced Encryption Standard (AES)** Daemen, J., & Rijmen, V. (2002). "The Design of Rijndael: AES - The Advanced Encryption Standard." *Springer Science & Business Media*.
The mathematical specification for the AES-256 cipher utilized by the underlying SQL-Cipher engine.
- 6. Vulnerability Scoring Metrics (CVSS & EPSS)** FIRST (Forum of Incident Response and Security Teams). (2019). "Common Vulnerability Scoring System v3.1: Specification Document."
Jacobs, J., et al. (2021). "Exploit Prediction Scoring System (EPSS)." *arXiv preprint arXiv:2108.11803*.
The standardized industry models integrated into the tools/risk_scorer.py module to calculate probabilistic threat vectors.

9. System Integrity and Self-Healing Architecture

Maintaining the stability of an orchestration engine executing over 50 third-party binaries is a continuous challenge. Binaries are frequently updated, dependency structures shift, and raw scanner outputs are inherently unpredictable. This chapter details the self-healing and system integrity mechanisms built directly into the SMP architecture to guarantee operational resilience.

9.1 The Pre-Flight System Checker (`system_checker.py`)

Prior to launching any orchestration workflows, SMP executes a rigorous pre-flight diagnostic routine via the `tools/system_checker.py` module. This checker operates as an autonomous gatekeeper, preventing the platform from launching into a corrupted state.

9.1.1 Cryptographic Binary Verification

Unlike traditional platforms that blindly trust binaries located within the system PATH, SMP verifies the cryptographic integrity of its underlying tools. During installation, `setup.sh` records the SHA-256 hash of every downloaded binary (e.g., `nuclei`, `httpx`, `ffuf`).

The `system_checker.py` module dynamically computes the SHA-256 hash of the binaries present in the `bin/` directory and compares them against the known-good signature matrix. - If a binary has been tampered with by a malicious actor (e.g., an attacker replacing `nmap` with a reverse-shell payload), the checker flags a CRITICAL failure and aborts the application startup. - If a binary is missing, the checker logs a WARNING, allowing the platform to degrade gracefully (as the DAG will dynamically bypass dependent tools).

9.1.2 Database Schema Validation

Before attempting to write pentest telemetry, the checker connects to `security.db` and performs a PRAGMA integrity check. It verifies that all required tables (`findings`, `targets`, `system_secrets`) exist and contain the correct column definitions. If an older database version is detected, the checker automatically executes `safe ALTER TABLE`

SQL migrations (e.g., retroactively adding the `centrality_score` column for the V9 Neural Brain update).

9.2 The Static Pipeline Verifier (`verify_smp.py`)

While `system_checker.py` validates the runtime environment, `verify_smp.py` serves as a static analysis tool for the codebase itself, heavily utilized within the Continuous Integration (CI) pipeline.

9.2.1 Graph Acyclicity and Deadlock Prevention

The most critical function of `verify_smp.py` is validating the Directed Acyclic Graph (DAG) established by the `scanners/` directory.

The script dynamically imports the `@register_scanner` decorators from all 57 scanner modules and constructs a virtual graph. It then executes a Depth-First Search (DFS) algorithm to mathematically prove the absence of cycles.

If a developer accidentally creates a cyclic dependency (e.g., Tool A depends on Tool B, which depends on Tool C, which depends on Tool A), the verifier immediately fails the CI pipeline, preventing a catastrophic infinite deadlock from reaching production.

```
# CI Verification Pipeline Trace
$ python3 tools/verify_smp.py
[INFO] Parsed 57 scanner modules successfully.
[INFO] Constructing Directed Acyclic Graph...
[SUCCESS] DFS validation passed: 0 cycles detected.
[SUCCESS] No orphaned dependencies.
[SUCCESS] DAG integrity mathematically verified.
```

9.3 Noise Reduction: The Levenshtein Deduplicator

A systemic flaw in executing 50 overlapping security tools is the massive generation of duplicate findings. For example, `sqlmap`, `wapiti`, and `nuclei` may all independently discover the same SQL Injection vulnerability on the same URL parameter.

Displaying three identical alerts induces extreme cognitive overload for the analyst. To resolve this, SMP employs a deterministic heuristic engine within `tools/finding_deduplicator.py`.

9.3.1 Levenshtein Distance Fuzzy Matching

Because different tools describe the exact same vulnerability using disparate terminology (e.g., “SQLi” vs “SQL Injection” vs “Blind SQL”), exact string matching is insufficient.

The deduplicator utilizes the **Levenshtein Distance** algorithm to calculate the mathematical edit distance between the titles and descriptions of findings affecting the same target endpoint.

$$\text{Similarity} = 1.0 - \left(\frac{\text{Levenshtein}(S_1, S_2)}{\max(|S_1|, |S_2|)} \right)$$

Findings that achieve a similarity ratio ≥ 0.82 are mathematically proven to be identical vulnerabilities. The engine automatically merges these findings, escalating the Confidence Score, and consolidating the visual representation within the Neural Brain. This process operates entirely autonomously in the background, drastically reducing the noise-to-signal ratio of the final PDF report.

Appendix A: Comprehensive Scanner Compendium

This appendix provides an exhaustive technical breakdown of every security scanner integrated into the Security Management Platform. The data presented herein is mathematically derived directly from the runtime registration metadata within the scanners/ directory.

Amass (amass.py)

Execution Step: Running Amass

Underlying Binary: amass

DAG Dependencies: [Subfinder]

Baseline Confidence Score: 85/100

Architectural Description

The `amass.py` module serves as the primary execution wrapper for the Amass tool. Because this tool relies on Subfinder, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `amass` is not present in the system PATH or the `bin/` directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Amass is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

API Fuzzer (api_fuzzer.py)

Execution Step: Running API Fuzzer

Underlying Binary: `**DAG Dependencies**:[Katana]`

Baseline Confidence Score: 80/100

Architectural Description

The `api_fuzzer.py` module serves as the primary execution wrapper for the API Fuzzer tool. Because this tool relies on Katana, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of API Fuzzer is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

Arjun (arjun.py)

Execution Step: Running Arjun

Underlying Binary: arjun

DAG Dependencies: [Dalfox]

Baseline Confidence Score: 85/100

Architectural Description

The arjun.py module serves as the primary execution wrapper for the Arjun tool. Because this tool relies on Dalfox, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If arjun is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Arjun is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Cloud Enum (cloud_enum.py)

Execution Step: Running Cloud Enum

Underlying Binary: cloud_enum

DAG Dependencies: [ParamSpider]

Baseline Confidence Score: 85/100

Architectural Description

The cloud_enum.py module serves as the primary execution wrapper for the Cloud Enum tool. Because this tool relies on ParamSpider, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If cloud_enum is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Cloud Enum is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

CMS Scanner (`cms_scanner.py`)

Execution Step: Running CMS Scanner

Underlying Binary: `‘**DAG Dependencies**:[CORS]’`

Baseline Confidence Score: 85/100

Architectural Description

The `cms_scanner.py` module serves as the primary execution wrapper for the CMS Scanner tool. Because this tool relies on CORS, the Kahn’s Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `‘is` is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`’.

DAG Memory Profiling

Upon execution, the standard output of CMS Scanner is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Commix (commix.py)

Execution Step: Running Commix

Underlying Binary: commix

DAG Dependencies: [Katana]

Baseline Confidence Score: 95/100

Architectural Description

The commix.py module serves as the primary execution wrapper for the Commix tool. Because this tool relies on Katana, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If commix is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Commix is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

CORS (cors_scanner.py)

Execution Step: Running CORS

Underlying Binary: `**DAG Dependencies**:[Robots.txt]`

Baseline Confidence Score: 95/100

Architectural Description

The `cors_scanner.py` module serves as the primary execution wrapper for the CORS tool. Because this tool relies on `Robots.txt`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`'.

DAG Memory Profiling

Upon execution, the standard output of CORS is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

CRLF Scanner (crlf_scanner.py)

Execution Step: Running CRLF Scanner

Underlying Binary: '**DAG Dependencies**:[Tech Fingerprint]'

Baseline Confidence Score: 85/100

Architectural Description

The `crlf_scanner.py` module serves as the primary execution wrapper for the CRLF Scanner tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of CRLF Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

CRT.sh (crtsh.py)

Execution Step: Running CRT.sh

Underlying Binary: `**DAG Dependencies**:[theHarvester]`

Baseline Confidence Score: 95/100

Architectural Description

The `crtsh.py` module serves as the primary execution wrapper for the `CRT.sh` tool. Because this tool relies on `theHarvester`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a `0` status code. If `'is not present in the systemPATH or the bin/directory'`, the orchestrator will automatically mark this node as `SKIPPED`.

DAG Memory Profiling

Upon execution, the standard output of `CRT.sh` is intercepted by the `Subprocess-Watchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Dalfox (dalfox.py)

Execution Step: Running Dalfox

Underlying Binary: dalfox

DAG Dependencies: [Gitleaks]

Baseline Confidence Score: 90/100

Architectural Description

The dalfox.py module serves as the primary execution wrapper for the Dalfox tool. Because this tool relies on Gitleaks, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If dalfox is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Dalfox is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

DNSx (dnsx.py)

Execution Step: Running DNSx

Underlying Binary: dnsx

DAG Dependencies: [Subfinder]

Baseline Confidence Score: 95/100

Architectural Description

The dnsx.py module serves as the primary execution wrapper for the DNSx tool. Because this tool relies on Subfinder, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If dnsx is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of DNSx is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Feroxbuster (feroxbuster.py)

Execution Step: Running Feroxbuster

Underlying Binary: feroxbuster

DAG Dependencies: [Katana]

Baseline Confidence Score: 85/100

Architectural Description

The feroxbuster.py module serves as the primary execution wrapper for the Feroxbuster tool. Because this tool relies on Katana, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If feroxbuster is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Feroxbuster is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

ffuf (ffuf.py)

Execution Step: Running ffuf

Underlying Binary: ffuf

DAG Dependencies: [Nuclei]

Baseline Confidence Score: 90/100

Architectural Description

The ffuf.py module serves as the primary execution wrapper for the ffuf tool. Because this tool relies on Nuclei, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If ffuf is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of ffuf is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Gitleaks (gitleaks.py)

Execution Step: Running Gitleaks

Underlying Binary: `**DAG Dependencies**:[Shodan]`

Baseline Confidence Score: 95/100

Architectural Description

The gitleaks.py module serves as the primary execution wrapper for the Gitleaks tool. Because this tool relies on Shodan, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If it is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Gitleaks is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

GraphQL Scanner (graphql_scanner.py)

Execution Step: Running GraphQL Scanner

Underlying Binary: '**DAG Dependencies**:[Katana]'

Baseline Confidence Score: 80/100

Architectural Description

The graphql_scanner.py module serves as the primary execution wrapper for the GraphQL Scanner tool. Because this tool relies on Katana, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of GraphQL Scanner is intercepted by the Sub-process Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

HackerTarget (hackertarget.py)

Execution Step: Running HackerTarget

Underlying Binary: `'**DAG Dependencies**:[CRT.sh]'`

Baseline Confidence Score: 90/100

Architectural Description

The `hackertarget.py` module serves as the primary execution wrapper for the HackerTarget tool. Because this tool relies on `CRT.sh`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `'is` is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`'.

DAG Memory Profiling

Upon execution, the standard output of HackerTarget is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Security Headers (headers_scanner.py)

Execution Step: Running Security Headers

Underlying Binary: '**DAG Dependencies**:[SSL]'

Baseline Confidence Score: 95/100

Architectural Description

The headers_scanner.py module serves as the primary execution wrapper for the Security Headers tool. Because this tool relies on SSL, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Security Headers is intercepted by the Sub-process Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

HTTPx (httpx_scanner.py)

Execution Step: Running HTTPx

Underlying Binary: httpx

DAG Dependencies: []

Baseline Confidence Score: 95/100

Architectural Description

The `httpx_scanner.py` module serves as the primary execution wrapper for the HTTPx tool. Because this tool relies on `'`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a `status` code. If `httpx` is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`.

DAG Memory Profiling

Upon execution, the standard output of HTTPx is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Auth Brute-Force Test (hydra_scanner.py)

Execution Step: Running Auth Brute-Force Test

Underlying Binary: `'**DAG Dependencies**:[Tech Fingerprint]'`

Baseline Confidence Score: 85/100

Architectural Description

The `hydra_scanner.py` module serves as the primary execution wrapper for the Auth Brute-Force Test tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATHor thebin/directory, the orchestrator will automatically mark this node asSKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Auth Brute-Force Test is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

JWT Scanner (jwt_scanner.py)

Execution Step: Running JWT Scanner

Underlying Binary: jwt_tool

DAG Dependencies: [Commix]

Baseline Confidence Score: 85/100

Architectural Description

The jwt_scanner.py module serves as the primary execution wrapper for the JWT Scanner tool. Because this tool relies on Commix, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If jwt_tool is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of JWT Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Katana (katana.py)

Execution Step: Running Katana

Underlying Binary: katana

DAG Dependencies: [HTTPx]

Baseline Confidence Score: 90/100

Architectural Description

The katana.py module serves as the primary execution wrapper for the Katana tool. Because this tool relies on HTTPx, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If katana is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Katana is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Masscan (masscan.py)

Execution Step: Running Masscan

Underlying Binary: masscan

DAG Dependencies: [WPScan]

Baseline Confidence Score: 95/100

Architectural Description

The masscan.py module serves as the primary execution wrapper for the Masscan tool. Because this tool relies on WPScan, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If masscan is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Masscan is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Nikto (nikto.py)

Execution Step: Running Nikto

Underlying Binary: nikto

DAG Dependencies: [CMS Scanner]

Baseline Confidence Score: 90/100

Architectural Description

The nikto.py module serves as the primary execution wrapper for the Nikto tool. Because this tool relies on CMS Scanner, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If nikto is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Nikto is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Nmap (nmap.py)

Execution Step: Running Nmap

Underlying Binary: nmap

DAG Dependencies: [Traceroute]

Baseline Confidence Score: 95/100

Architectural Description

The `nmap.py` module serves as the primary execution wrapper for the Nmap tool. Because this tool relies on Traceroute, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `nmap` is not present in the system `PATH` or the `bin/` directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Nmap is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Nuclei (nuclei.py)

Execution Step: Running Nuclei

Underlying Binary: nuclei

DAG Dependencies: [Nikto]

Baseline Confidence Score: 95/100

Architectural Description

The nuclei.py module serves as the primary execution wrapper for the Nuclei tool. Because this tool relies on Nikto, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If nuclei is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Nuclei is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Open Redirect (open_redirect.py)

Execution Step: Running Open Redirect

Underlying Binary: `**DAG Dependencies**:[ffuf]`

Baseline Confidence Score: 90/100

Architectural Description

The `open_redirect.py` module serves as the primary execution wrapper for the Open Redirect tool. Because this tool relies on `ffuf`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `'is'` is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`.

DAG Memory Profiling

Upon execution, the standard output of Open Redirect is intercepted by the `SubprocessWatchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

ParamSpider (paramspider.py)

Execution Step: Running ParamSpider

Underlying Binary: paramspider

DAG Dependencies: [Masscan]

Baseline Confidence Score: 85/100

Architectural Description

The paramspider.py module serves as the primary execution wrapper for the ParamSpider tool. Because this tool relies on Masscan, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If paramspider is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of ParamSpider is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Path Traversal Scanner (path_traversal.py)

Execution Step: Running Path Traversal Scanner

Underlying Binary: '**DAG Dependencies**:[Tech Fingerprint]'

Baseline Confidence Score: 85/100

Architectural Description

The path_traversal.py module serves as the primary execution wrapper for the Path Traversal Scanner tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATHor thebin/directory, the orchestrator will automatically mark this node asSKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Path Traversal Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Retire.js Scanner (retire_js.py)

Execution Step: Running Retire.js Scanner

Underlying Binary: '**DAG Dependencies**:[Tech Fingerprint]'

Baseline Confidence Score: 80/100

Architectural Description

The retire_js.py module serves as the primary execution wrapper for the Retire.js Scanner tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Retire.js Scanner is intercepted by the Sub-process Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

Robots.txt (robots_scanner.py)

Execution Step: Running Robots.txt

Underlying Binary: `**DAG Dependencies**:[Security Headers]`

Baseline Confidence Score: 95/100

Architectural Description

The `robots_scanner.py` module serves as the primary execution wrapper for the `Robots.txt` tool. Because this tool relies on Security Headers, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of `Robots.txt` is intercepted by the `Subprocess-Watchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

Shodan (shodan_idb.py)

Execution Step: Running Shodan

Underlying Binary: `**DAG Dependencies**:[SQLMap]`

Baseline Confidence Score: 90/100

Architectural Description

The `shodan_idb.py` module serves as the primary execution wrapper for the Shodan tool. Because this tool relies on SQLMap, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If it is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`.

DAG Memory Profiling

Upon execution, the standard output of Shodan is intercepted by the `Subprocess-Watchdog`. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

HTTP Smuggling Scanner (smuggler.py)

Execution Step: Running HTTP Smuggling Scanner

Underlying Binary: '**DAG Dependencies**:[Nmap]'

Baseline Confidence Score: 80/100

Architectural Description

The smuggler.py module serves as the primary execution wrapper for the HTTP Smuggling Scanner tool. Because this tool relies on Nmap, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of HTTP Smuggling Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

SQLMap (sqlmap.py)

Execution Step: Running SQLMap

Underlying Binary: sqlmap

DAG Dependencies: [Wapiti]

Baseline Confidence Score: 95/100

Architectural Description

The sqlmap.py module serves as the primary execution wrapper for the SQLMap tool. Because this tool relies on Wapiti, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If sqlmap is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of SQLMap is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

SSL (ssl_scanner.py)

Execution Step: Running SSL Scan

Underlying Binary: `**DAG Dependencies**:[Nmap]`

Baseline Confidence Score: 95/100

Architectural Description

The `ssl_scanner.py` module serves as the primary execution wrapper for the SSL tool. Because this tool relies on Nmap, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of SSL is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

SSRF Scanner (ssrf_scanner.py)

Execution Step: Running SSRF Scanner

Underlying Binary: '**DAG Dependencies**:[Tech Fingerprint]'

Baseline Confidence Score: 85/100

Architectural Description

The `ssrf_scanner.py` module serves as the primary execution wrapper for the SSRF Scanner tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of SSRF Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Subfinder (subfinder.py)

Execution Step: Running Subfinder

Underlying Binary: subfinder

DAG Dependencies: []

Baseline Confidence Score: 90/100

Architectural Description

The subfinder.py module serves as the primary execution wrapper for the Subfinder tool. Because this tool relies on `'`, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with `a0status` code. If subfinder is not present in the `systemPATH` or the `bin/directory`, the orchestrator will automatically mark this node as `SKIPPED`.

DAG Memory Profiling

Upon execution, the standard output of Subfinder is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Tech Fingerprint (tech_fingerprint.py)

Execution Step: Running Tech Fingerprint

Underlying Binary: '**DAG Dependencies**:[Open Redirect]'

Baseline Confidence Score: 85/100

Architectural Description

The tech_fingerprint.py module serves as the primary execution wrapper for the Tech Fingerprint tool. Because this tool relies on Open Redirect, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or thebin/directory, the orchestrator will automatically mark this node asSKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Tech Fingerprint is intercepted by the Sub-processWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

theHarvester (theharvester.py)

Execution Step: Running theHarvester

Underlying Binary: `**DAG Dependencies**:[Subfinder]`

Baseline Confidence Score: 80/100

Architectural Description

The theharvester.py module serves as the primary execution wrapper for the theHarvester tool. Because this tool relies on Subfinder, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of theHarvester is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

Traceroute (traceroute.py)

Execution Step: Running Traceroute

Underlying Binary: traceroute

DAG Dependencies: [Wayback Machine]

Baseline Confidence Score: 90/100

Architectural Description

The traceroute.py module serves as the primary execution wrapper for the Traceroute tool. Because this tool relies on Wayback Machine, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If traceroute is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Traceroute is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Wapiti (wapiti.py)

Execution Step: Running Wapiti

Underlying Binary: wapiti

DAG Dependencies: [Tech Fingerprint]

Baseline Confidence Score: 90/100

Architectural Description

The wapiti.py module serves as the primary execution wrapper for the Wapiti tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If wapiti is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Wapiti is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

Wayback Machine (wayback.py)

Execution Step: Running Wayback Machine

Underlying Binary: `**DAG Dependencies**:[Whois]`

Baseline Confidence Score: 80/100

Architectural Description

The wayback.py module serves as the primary execution wrapper for the Wayback Machine tool. Because this tool relies on Whois, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is' is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of Wayback Machine is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 80.

WhatWeb (whatweb.py)

Execution Step: Running WhatWeb

Underlying Binary: whatweb

DAG Dependencies: [HTTPx]

Baseline Confidence Score: 85/100

Architectural Description

The `whatweb.py` module serves as the primary execution wrapper for the WhatWeb tool. Because this tool relies on HTTPx, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If `whatweb` is not present in the system PATH or the `bin/` directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of WhatWeb is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the `dag_state` SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Whois (whois_scanner.py)

Execution Step: Running Whois

Underlying Binary: whois

DAG Dependencies: [HackerTarget]

Baseline Confidence Score: 95/100

Architectural Description

The whois_scanner.py module serves as the primary execution wrapper for the Whois tool. Because this tool relies on HackerTarget, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If whois is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of Whois is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 95.

WPScan (wpscan.py)

Execution Step: Running WPScan

Underlying Binary: wpscan

DAG Dependencies: [JWT Scanner]

Baseline Confidence Score: 90/100

Architectural Description

The wpscan.py module serves as the primary execution wrapper for the WPScan tool. Because this tool relies on JWT Scanner, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If wpscan is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of WPScan is intercepted by the Subprocess-Watchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 90.

XXE Scanner (xxe_scanner.py)

Execution Step: Running XXE Scanner

Underlying Binary: '**DAG Dependencies**:[Tech Fingerprint]'

Baseline Confidence Score: 85/100

Architectural Description

The xxe_scanner.py module serves as the primary execution wrapper for the XXE Scanner tool. Because this tool relies on Tech Fingerprint, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If 'is not present in the systemPATH or the bin/directory, the orchestrator will automatically mark this node as SKIPPED'.

DAG Memory Profiling

Upon execution, the standard output of XXE Scanner is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

ZAP (zap.py)

Execution Step: Running ZAP

Underlying Binary: zap

DAG Dependencies: [Cloud Enum]

Baseline Confidence Score: 85/100

Architectural Description

The zap.py module serves as the primary execution wrapper for the ZAP tool. Because this tool relies on Cloud Enum, the Kahn's Topological Sort algorithm explicitly prevents its execution until those dependencies successfully exit with a 0 status code. If zap is not present in the system PATH or the bin/ directory, the orchestrator will automatically mark this node as SKIPPED.

DAG Memory Profiling

Upon execution, the standard output of ZAP is intercepted by the SubprocessWatchdog. The execution thread calculates the delta between the start epoch and end epoch, updating the dag_state SQLite table in real-time. Results are piped into the Neural Brain matrix, weighted by the base confidence score of 85.

Appendix B: Database Schemas and Data Dictionaries

To ensure localized data sovereignty and high-performance querying, the Security Management Platform (SMP) persists state across three discrete SQLite databases. This appendix documents the formal Data Definition Language (DDL) and schema architecture utilized in V9.4.0.

B.1 The Encrypted Pentest Database (security.db)

This database contains all highly sensitive topological and vulnerability data collected during a VAPT engagement. It is encrypted at rest using SQLCipher (AES-256 in CBC mode) with a PBKDF2 HMAC-SHA256 derived key (600,000 iterations).

targets Table

Stores the primary domain or IP scope parameters for an engagement.

```
CREATE TABLE targets (  
    target_id INTEGER PRIMARY KEY AUTOINCREMENT,  
    url TEXT NOT NULL UNIQUE,  
    added_date DATETIME DEFAULT CURRENT_TIMESTAMP,  
    last_scanned DATETIME,  
    attestation_signed BOOLEAN DEFAULT 0,  
    scan_profile TEXT DEFAULT 'standard'  
);
```

findings Table

The core relational table mapping vulnerability telemetry to the target scope.

```
CREATE TABLE findings (  
    finding_id INTEGER PRIMARY KEY AUTOINCREMENT,  
    scan_id INTEGER NOT NULL,  
    tool TEXT NOT NULL,  
    severity TEXT NOT NULL,  
    title TEXT NOT NULL,  
    description TEXT,
```

```

    confidence INTEGER DEFAULT 50,
    cve_id TEXT,
    epss_score REAL,
    centrality_score REAL DEFAULT 0.0,
    FOREIGN KEY(scan_id) REFERENCES scans(scan_id) ON DELETE CASCADE
);

```

system_secrets Table

Stores the internally generated cryptographic keys (e.g., Fernet keys) required to decrypt the unstructured evidence blobs stored on the host filesystem.

```

CREATE TABLE system_secrets (
    secret_id INTEGER PRIMARY KEY AUTOINCREMENT,
    key_name TEXT NOT NULL UNIQUE,
    key_value TEXT NOT NULL, -- The Base64 encoded Fernet Key
    created_at DATETIME DEFAULT CURRENT_TIMESTAMP
);

```

B.2 The Public Intelligence Database (global_intel.db)

This database powers the “Neural Brain” heuristic engine. Because it contains exclusively public, mathematical models (and zero client-specific data), it is intentionally unencrypted to maximize concurrent SELECT query performance.

cve_cache Table

A localized cache of the National Vulnerability Database (NVD) to prevent redundant outbound API calls and rate-limiting.

```

CREATE TABLE cve_cache (
    cve_id TEXT PRIMARY KEY,
    cvss_v3_score REAL,
    cvss_vector TEXT,
    description TEXT,
    published_date DATETIME,
    last_modified DATETIME
);

```

epss_metrics Table

The Exploit Prediction Scoring System parameters, updated daily.

```

CREATE TABLE epss_metrics (
    cve_id TEXT PRIMARY KEY,
    epss_probability REAL NOT NULL, -- Range: 0.0 to 1.0
    percentile REAL,
    date_fetched DATETIME DEFAULT CURRENT_TIMESTAMP
);

```

cisa_kev Table

The US Cybersecurity and Infrastructure Security Agency's Known Exploited Vulnerabilities catalog.

```
CREATE TABLE cisa_kev (
    cve_id TEXT PRIMARY KEY,
    vendor_project TEXT,
    product TEXT,
    vulnerability_name TEXT,
    date_added DATETIME,
    due_date DATETIME,
    known_ransomware_campaign_use TEXT
);
```

B.3 The Operational Redundancy Database (redundancy.db)

Also encrypted via SQLCipher, this database acts as a localized transaction log to recover state in the event of a catastrophic system failure (e.g., power loss during a 6-hour scan).

dag_state Table

Tracks the topological sorting queue and completion status of the current scan.

```
CREATE TABLE dag_state (
    scan_id INTEGER NOT NULL,
    node_name TEXT NOT NULL,
    in_degree INTEGER NOT NULL,
    status TEXT DEFAULT 'PENDING', -- PENDING, RUNNING, COMPLETED, FAILED
    start_time DATETIME,
    end_time DATETIME,
    PRIMARY KEY (scan_id, node_name)
);
```

blob_pointers Table

Maintains the mapping between relational findings_id and the encrypted Fernet flat-files residing in reports/evidence/.

```
CREATE TABLE blob_pointers (
    pointer_id INTEGER PRIMARY KEY AUTOINCREMENT,
    finding_id INTEGER NOT NULL,
    file_path TEXT NOT NULL,
    file_hash TEXT NOT NULL, -- SHA-256 for integrity verification
    FOREIGN KEY(finding_id) REFERENCES findings(finding_id) ON DELETE CASCADE
);
```

Appendix C: Comprehensive REST API Documentation

The Security Management Platform can operate in a completely headless state via its FastAPI backend. This allows enterprise integration into existing CI/CD pipelines (e.g., Jenkins, GitLab CI). This appendix serves as the definitive reference manual for all available REST endpoints.

C.1 Authentication and Authorization

By default, the REST API enforces strict Bearer Token authentication. All requests to protected endpoints must include the Authorization header.

POST /api/v1/auth/login

Generates a short-lived JSON Web Token (JWT) for API access.

Request Body:

```
{
  "username": "admin",
  "password": "SuperSecretMasterPassword"
}
```

Response (200 OK):

```
{
  "access_token": "eyJhbGciOiJIUzI1NiIsInR...\"",
  "token_type": "bearer",
  "expires_in": 3600
}
```

C.2 Orchestration Endpoints

These endpoints control the Directed Acyclic Graph (DAG) orchestration engine.

POST /api/v1/scans/start

Initiates a new VAPT engagement against a target scope.

Request Body:

```
{
  "target": "https://example.com",
  "profile": "aggressive",
  "exclude_tools": ["amass", "sqlmap"],
  "attestation_signed": true
}
```

Response (202 Accepted):

```
{
  "scan_id": 42,
  "status": "INITIALIZED",
  "estimated_completion_sec": 7200
}
```

GET /api/v1/scans/{scan_id}/status

Retrieves the real-time execution state of the topological DAG.

Response (200 OK):

```
{
  "scan_id": 42,
  "global_status": "RUNNING",
  "progress_percentage": 45.2,
  "active_nodes": ["nuclei", "ffuf", "dalfox"],
  "completed_nodes": ["nmap", "httpx", "subfinder"],
  "failed_nodes": []
}
```

POST /api/v1/scans/{scan_id}/abort

Sends a POSIX SIGKILL signal to all running child processes associated with the scan ID and safely unwinds the SQLite transaction logs.

Response (200 OK):

```
{
  "scan_id": 42,
  "status": "ABORTED",
  "processes_terminated": 3
}
```

C.3 Intelligence and Telemetry Endpoints

These endpoints retrieve the mathematically correlated vulnerabilities generated by the Neural Brain.

GET /api/v1/findings/{scan_id}

Retrieves all raw findings associated with a specific engagement.

Query Parameters: - min_severity (optional): low, medium, high, critical. - tool_filter (optional): e.g., nuclei.

Response (200 OK):

```
{
  "count": 142,
  "findings": [
    {
      "finding_id": 1024,
      "tool": "nuclei",
      "severity": "high",
      "title": "Exposed .git Repository",
      "epss_score": 0.84,
      "centrality": 0.12
    }
  ]
}
```

GET /api/v1/intelligence/attack_chains

Retrieves the clustered semantic attack chains generated by the TF-IDF cosine similarity matrix.

Response (200 OK):

```
{
  "clusters_found": 3,
  "clusters": [
    {
      "cluster_id": "c_9f8a",
      "primary_vector": "Cross-Site Scripting (Reflected)",
      "associated_finding_ids": [14, 88, 92],
      "confidence": 98.5
    }
  ]
}
```

GET /api/v1/intelligence/linchpins

Retrieves the topological chokepoints identified by the Degree Centrality algorithm.

Response (200 OK):

```
{
  "linchpins": [
    {
      "component": "api.internal.local",
      "centrality_score": 0.98,
      "critical_cves": 4
    }
  ]
}
```

C.4 Data Sovereignty & Evidence

GET /api/v1/evidence/{finding_id}/download

Retrieves the raw, unstructured evidence blob (e.g., full HTTP request/response traces).

Note: The API automatically decrypts the Fernet blob in memory using the Master Key prior to transmitting it over the TLS socket.

Response (200 OK, Content-Type: application/json):

```
{
  "finding_id": 1024,
  "raw_output": "GET /.git/config HTTP/1.1\nHost: example.com\n\nHTTP/1.1 200 OK\n\n[co
}"
```

10. Glossary of Terms

AES-256 (Advanced Encryption Standard)

A symmetric block cipher utilized by the U.S. government to protect classified information. SMP uses the 256-bit key length variant within SQLCipher.

CISA KEV (Known Exploited Vulnerabilities)

A definitive catalog maintained by the Cybersecurity and Infrastructure Security Agency listing CVEs actively used in cyber attacks.

Cosine Similarity

A mathematical measure of similarity between two non-zero vectors. SMP uses this within the TF-IDF clustering engine to mathematically group related vulnerabilities.

CVE (Common Vulnerabilities and Exposures)

A standardized dictionary of publicly known information security vulnerabilities and exposures.

CVSS (Common Vulnerability Scoring System)

An open industry standard for assessing the severity of computer system security vulnerabilities.

DAG (Directed Acyclic Graph)

A mathematical graph structure that flows in one direction and contains no cycles. Used in SMP to manage non-linear orchestration dependencies.

EPSS (Exploit Prediction Scoring System)

A data-driven model for estimating the likelihood (probability) that a software vulnerability will be exploited in the wild.

Fernet

A symmetric encryption specification utilizing AES-128 in CBC mode, PKCS7 padding, and HMAC-SHA256 for authentication. Used in SMP for encrypting raw tool output blobs on disk.

Kahn's Algorithm

An algorithm used to find a topological ordering of a directed acyclic graph. SMP uses this to compute execution order based on in-degree dependencies.

Levenshtein Distance

A string metric for measuring the difference between two sequences. SMP uses this to deduplicate findings.

Local-First

A software architecture paradigm emphasizing that the primary copy of data should reside on the local device, rather than on a remote cloud server, ensuring maximum privacy and sovereignty.

PBKDF2 (Password-Based Key Derivation Function 2)

A cryptographic algorithm that derives a strong, fixed-length key from a variable-length password to prevent brute-force dictionary attacks. SMP strictly enforces 600,000 iterations.

PySide6

The official Python bindings for the Qt framework, providing access to native C++ GUI components.

SQLCipher

An open-source extension to SQLite that provides transparent 256-bit AES encryption of database files.

TF-IDF (Term Frequency-Inverse Document Frequency)

A numerical statistic intended to reflect how important a word is to a document in a collection or corpus. Used by the Neural Brain for semantic clustering.

11. Index

A

AES-256, 5.1 Air-gapped, 1.1, 7.3 API (FastAPI), 3.2

C

CISA KEV, B.2 Cosine Similarity, 4.2 Cryptography, Chapter 5 CVSS (Common Vulnerability Scoring System), B.2

D

DAG (Directed Acyclic Graph), 4.1 Data Sovereignty, 1.1, 5.0 Degree Centrality, 4.3 Docker, 2.2

E

EPSS (Exploit Prediction Scoring System), B.2 EventBus, 3.4

F

Fernet Encryption, 5.2

H

HMAC-SHA256, 5.1.1, 5.2.1

K

Kahn's Algorithm, 4.1.1 Key Derivation (PBKDF2), 5.1.1

L

Levenshtein Distance, 9.3

N

Nmap, 1.1, A.1 Neural Brain, 4.2, 4.3, 6.2 NVD (National Vulnerability Database), B.2

P

PageRank, 4.3 PBKDF2, 5.1.1 ProcessPoolExecutor, 4.1.2 PySide6, 3.1

S

Scanner Registry, 3.3 SQLCipher, 5.1, B.1 SQLite, 5.1 Subprocess Watchdog, A.1
System Checker, 9.1

T

TF-IDF, 4.2 Topological Sorting, 4.1.1